LinkedIn Thought Leadership Agent

An agentic AI skills pack that transforms a single strategic intent into a 6-week LinkedIn content series — complete with narrative arcs, voice consistency, human-in-the-loop checkpoints, and automated posting via MCP.

> ***MSIS 549 B — Machine Learning & AI for Business Applications*** *University of Washington, Winter 2025-2026 | Homework 2: Agentic AI for Real-World Impact*

The Problem

Creating consistent, high-quality LinkedIn thought leadership is time-consuming and cognitively demanding. Most professionals default to one of two failure modes:

1. **Generic AI output** — "In today's data-driven world..." posts that all sound the same
2. **Sporadic posting** — starting strong, then going silent for weeks

This system solves both by decomposing content creation into specialized skills that enforce quality, consistency, and strategic coherence across a 6-week publishing cadence.

Architecture

The system follows **Path A** (Agentic Skills Pack) with **9 specialized skills** orchestrated by a master controller:

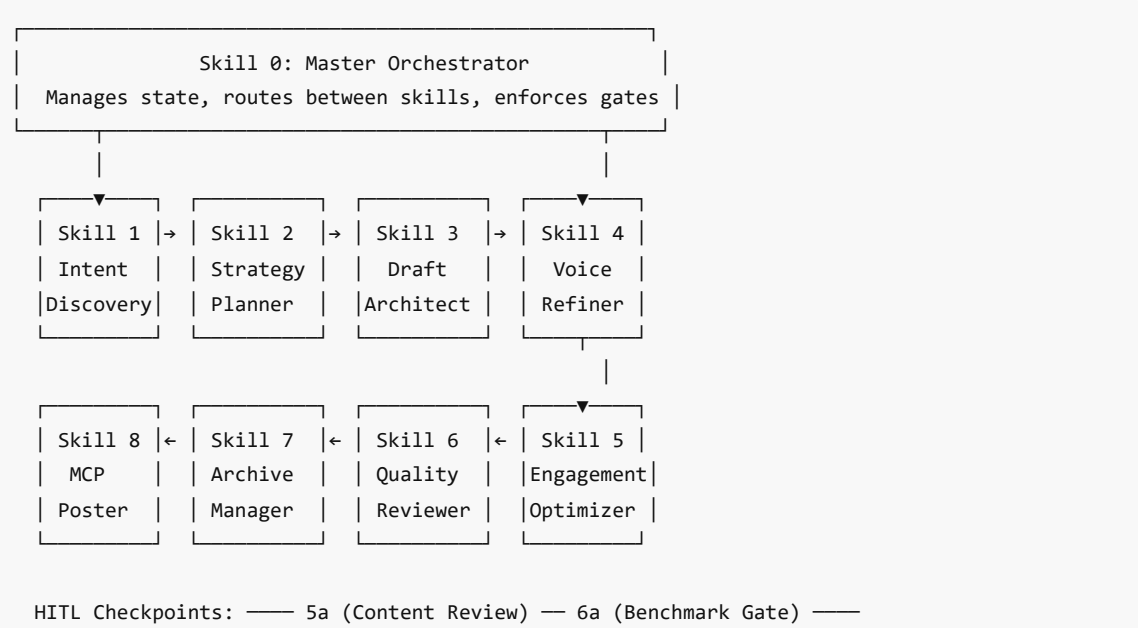

Skills Summary

#	Skill	Purpose
0	Master Orchestrator	Pipeline control, state management, error recovery
1	Intent Discovery	5-question interview to extract strategic intent
2	Content Strategist	6-week narrative arc with Hook→Framework→Story→Tactics→Vision→CTA structure

3	Draft Architect	First drafts with anti-AI-ism rules (no "In today's...", no "Let's dive in...")
4	Voice & Tone Refiner	Detects and replaces 15+ AI patterns with human voice
5	Engagement Optimizer	LinkedIn-specific optimization (hooks, CTAs, hashtags, visuals) + HITL 5a
6	Quality Reviewer	5-metric scoring rubric (1-5 scale) + HITL 6a benchmark gate
7	Archive Manager	Structured archival with full metadata for reproducibility
8	Poster & Reviewer	LiGo MCP integration for automated LinkedIn posting

Key Design Decisions

1. **Anti-AI-ism as a first-class concern** — Skills 3-4 contain explicit banned-pattern lists and replacement tables
2. **Two mandatory HITL checkpoints** — Content never publishes without human approval
3. **Narrative cohesion by design** — Skill 2 plans the full 6-week arc before any drafting begins
4. **MCP integration** — Skill 8 uses LiGo MCP for direct LinkedIn posting (with manual fallback)

Benchmark Results

Test Case	Actionability	Voice	Depth	Cohesion	LinkedIn	**Overall**
SQL Performance (Agentic)	4.2	4.7	4.3	4.8	4.3	**4.5**
Data Cleanliness (Agentic)	4.2	4.2	4.0	4.5	4.0	**4.2**
WAL Protocol (Edge Case)	3.5	2.5	4.0	3.8	3.2	**3.4**
Vague Input (Ambiguous)	3.5	3.8	3.2	4.0	3.5	**3.6**
Single-Prompt Baseline	**2.5**	**2.2**	**1.8**	**1.0**	**2.5**	**2.0**

Key finding: The agentic system outperformed the single-prompt baseline by **+2.5 points** on the primary test case (4.5 vs 2.0). Biggest improvement was in Narrative Cohesion (+3.8), which a single prompt fundamentally cannot achieve.

See `benchmark/BENCHMARK_APPENDIX.md` for full methodology, scoring rubrics, and failure analysis.

Repository Structure

```
├── README.md                          # This file
├── TUTORIAL_WRITEUP.md                # Full tutorial (assignment submission)
├── skills/
│   ├── skill_0_master_orchestrator.md # Pipeline orchestration
│   ├── skill_1_intent_discovery.md    # Strategic intent interview
│   ├── skill_2_content_strategist.md  # 6-week arc planning
│   ├── skill_3_draft_architect.md     # First draft generation
│   ├── skill_4_voice_tone_refiner.md  # AI pattern detection & replacement
│   ├── skill_5_engagement_optimizer.md# LinkedIn optimization + HITL 5a
│   ├── skill_6_quality_reviewer.md    # 5-metric scoring + HITL 6a
```

```
│   ├── skill_7_archive_manager.md     # Structured archival
│   └── skill_8_poster_reviewer.md     # LiGo MCP integration
├── benchmark/
│   └── BENCHMARK_APPENDIX.md          # Full benchmark methodology & results
├── outputs/
│   ├── Q1_2026_SQL_Performance_Series.md  # Generated 6-week series
│   └── 2026_Roadmap_Plan.md               # Annual content roadmap
├── diagrams/
│   └── architecture_infographic.html  # A4 landscape architecture diagram
└── archive/                           # Session archives (generated at runtime)
```

How to Use

Prerequisites

- Claude Code (or any Claude-based IDE with skill file support)
- Optional: LiGo MCP for automated LinkedIn posting

Quick Start

1. Clone this repository
2. Copy the `skills/` directory to your Claude Code skills folder (`~/.claude/skills/`)
3. Start a conversation and say: "*I want to create a LinkedIn thought leadership series about [your topic]*"
4. The orchestrator will guide you through:
    - **Intent Discovery** — 5-question interview to clarify your topic, audience, and message
    - **Content Strategy** — A 6-week narrative arc tailored to your intent
    - **Draft Generation** — 6 polished posts with anti-AI-ism enforcement
    - **Human Review** — Two checkpoints where you approve or revise content
    - **Publishing** — Direct to LinkedIn via MCP or manual copy/paste

Example Input

```
Topic: Optimizing SQL Query Performance
Audience: Data Engineers & Business Stakeholders
Core Message: Continuous improvement in SQL is essential for AI-readiness
Anecdote: Business team frustrated when data wasn't in sync with AI models
Tone: Provocative + Educational
```

Example Output

See `outputs/Q1_2026_SQL_Performance_Series.md` for the full 6-week series generated from the input above.

Tech Stack

- **LLM**: Claude (via Manus AI agentic platform)
- **MCP Integration**: LiGo MCP for LinkedIn API access
- **Evaluation**: Human rubric scoring (5-metric, 1-5 scale)
- **Architecture**: Path A — Agentic Skills Pack (markdown skill files)

Known Limitations

1. **Voice Consistency on technical topics** — Skill 4's voice refiner is optimized for Manager/Director audiences. Deeply technical audiences (e.g., database kernel engineers) score lower (2.5/5) because the

refiner over-simplifies jargon.
2. **Single-rater evaluation** — Benchmark scores are from one human evaluator. Inter-rater reliability with 2+ evaluators would strengthen the results.
3. **No A/B testing** — Posts haven't been tested against actual LinkedIn engagement metrics yet.

License

Author

Taashim Ahmed — University of Washington, MSIS 549 B (Winter 2025-2026)